

Matter-first automated discovery of transient spacetime shortcuts

Nikita M. Shirokov*

Independent Researcher, Moscow, Russia

(Dated: July 20, 2026)

We present an automated matter-first pipeline for discovering warp-precursor spacetimes—transient spacetime shortcuts: an optimizer proposes matter, a constraint solver (**GR_{Tresna}**) builds initial data, a GPU evolution code (**GR_{Teclyn}**) advances the full 3+1 Einstein equations, and null geodesics integrated through the *time-evolving* metric test whether light arrives earlier than in flat space. MAP-Elites quality-diversity search over a 39-dimensional bicomplex Q-ball genome, followed by CMA-ES refinement, yields a five-lump champion (candidate 146) supported by two canonical and three phantom (negative-energy) scalar sources. In the headline validation run ($L = 128$, $N = 256$, three refinement levels, $t = 30$), null rays arrive 20.2% earlier than the matched flat path, with all quality gates passed. The result is stable across a three-run resolution ladder ($N = 192/256/384$; fine–medium relative difference $\sim 0.1\%$) and a matched-spacing domain ladder ($L = 96$ and $L = 160$; within 0.12% of the baseline). An external ignition controller switches off at $t = 4$; because the geodesic that defines the quoted advantage launches at $t_{\text{emit}} = 4$, the timing difference is measured in freely evolving Einstein–Klein–Gordon dynamics, and a controller-free twin run reproduces it at 20.5%. A matched canonical-only (positive-energy) search finds no evolving timing advantage. In all cases the channel is transient: confinement decays as the supporting matter disperses.

I. INTRODUCTION

The goal of this work is an *automatic discovery pipeline* for faster-than-light (FTL) warp precursors, in the spirit of recent automated theory- and pipeline-discovery efforts in fundamental physics [1, 2]. We define a warp precursor—a transient *spacetime shortcut*—as a constraint-satisfying matter configuration that generates a transient, effective null timing advantage prior to its own dispersal: null rays between prescribed fixed endpoints in the weak-field region traverse the curved spacetime in less coordinate time than the matched flat-space path. Rather than hand-crafting a metric and hoping the implied stress-energy is physical, we want a closed loop that proposes candidates, evolves them under the Einstein equations, and retains only those that pass a gauge-invariant transport test.

The standard route is *geometry-first* (equivalently metric-first): one prescribes $g_{\mu\nu}$, reads off $T_{\mu\nu} = G_{\mu\nu}/8\pi$, and studies frozen slices [3–5]. That recipe is poorly suited to automated search. The implied matter is rarely the output of a consistent initial-value problem, and under dynamical evolution it typically disperses or collapses, destroying the intended channel. For a quality-diversity optimizer that must evaluate hundreds of candidates, a geometry-first genome is especially poorly conditioned: most deformations of $g_{\mu\nu}$ either violate the constraints, have no local matter realization, or fail as soon as time advances—so the archive fills with numerical failures rather than competing mechanisms.

We therefore invert the loop (Fig. 1). The optimizer proposes matter; **GR_{Tresna}** (a constraint solver that returns initial data satisfying the Hamiltonian and momentum constraints) builds the initial slice; **GR_{Teclyn}** [6, 7] (a

GPU-accelerated, AMReX-based CCZ4 evolution code) advances the spacetime; and probes ask whether a null shortcut appears. In this paper “FTL” and “null shortcut” mean the same operational claim: an *effective* reduction of null travel time relative to flat space between the same fixed weak-field endpoints—never local superluminal motion outside the light cone. The fractional advantage is f_{geo} [Eq. (13)]; for example $f_{\text{geo}} = 0.20$ means the curved transit is 20% shorter than the flat comparison. By classical theorems on superluminal travel and gravitational time delay [8–10], any such advantage between endpoints in the (near-)flat region requires violation of the null energy condition, and generic warp-drive spacetimes share this requirement [11]; the phantom sector introduced below is the controlled realization of that necessity in our matter model.

This criterion does real work in the pipeline: it is the acceptance test, not a post-hoc diagnostic. A warp precursor is useful only if it changes what light can do between fixed distant observers. Local coordinate signal speeds, gauge artifacts of the lapse and shift, or a positive frozen-slice diagnostic can look “superluminal” while no ray actually connects the endpoints sooner once the metric evolves. Asking whether null geodesics arrive earlier than in flat space is therefore the operational definition of the search target: a gauge-invariant transport certificate that the optimizer must earn, not a post-hoc plot.

Equally important is the *scoring metric* that turns each evolution into the fitness $S(\theta)$ used by the archive [Eq. (9)]. MAP-Elites converges to whatever S rewards. In earlier campaigns with weaker objectives the search reliably “cheated”: it harvested coordinate proxies while matter dispersed (“disperse-then-warp”), banked frozen-slice positives that failed the evolving geodesic test, or climbed scores through numerical pathologies—early horizons, integrator crashes, and other simulation-failure modes that accidentally produced large proxy signals.

* shirokov.nm@phystech.edu




FIG. 1. Geometry-first versus matter-first discovery. Dashed: prescribe a metric and infer its stress-energy. Solid: propose matter, solve and test an initial-value problem (IVP), evolve it, and archive only candidates that pass the 4D null-transport test.

The present objective therefore rewards only the evolving 4D shortcut and persistence, gates health terms by nontriviality, penalizes exotic excess and horizon formation, and dispersion-gates transport credit so a flying-apart cloud cannot buy a high score. Without that metric design the archive does not discover physics; it exploits weaknesses of the objective.

The search engine is MAP-Elites [12], a quality-diversity algorithm that keeps one elite per cell of a behavior archive rather than a single global optimum. We prefer this over a pure CMA-ES [13] hill-climb, which collapses onto one basin and discards diverse near-misses, and over reinforcement learning, which would require dense credit assignment across expensive GPU evolutions with sparse, delayed rewards. After MAP-Elites identifies a high-scoring basin, CMA-ES warm-starts from the archive champion and locally refines the genome before any high-resolution replay or resolution ladder (Sec. III). Both discovery stages run on a deliberately coarse grid chosen for evaluation throughput; production claims are deferred to the later high-quality (HQ) production replay and resolution matrix (Sec. II J).

We acknowledge that the headline champion (candidate 146) relies on a phantom scalar sector and therefore does *not* bypass the standard energy-condition violations associated with FTL spacetimes [8, 9, 11]. The value of this work is that an automated, matter-first Cauchy pipeline can discover, construct, and dynamically track constraint-satisfying FTL geometries—a necessary first step toward computationally searching the positive-energy sector with the same falsification stack. Claimed positive-energy warp solitons [14, 15] make that sector a live question, although generic-warp-drive analyses argue such constructions still violate the null energy condition [11]. As a first probe, a matched canonical-only (positive-energy) MAP-Elites search within the present ansatz yields no evolving timing advantage—every accepted evaluation returns $f_{\text{geo}}^{\text{evol}} = 0$ (Sec. VIB), so within this ansatz the shortcut requires the phantom sector.

A. Methodological Rigor

Three pillars distinguish this pipeline from geometry-first studies:

1. **Initial value problem:** constraint-clean data

from physical scalar sources, not coordinate scripts.

2. **Dynamic feedback:** full BSSN/CCZ4 evolution; dispersing matter collapses the channel.
3. **4D geodesic falsification:** the headline diagnostic is `ftl_geo_evolving`—null rays integrated through the *time-evolving* metric stack. Frozen-slice geodesics are retained as diagnostics but are not the transport certificate.

B. Contributions

- (i) An end-to-end matter-first discovery pipeline—MAP-Elites proposal, CMA-ES local refinement of the archive champion, constraint solve, GPU evolution, and matched evolving-geodesic scoring—with an explicit reward rule applied to every evaluation (Secs. II–III).
- (ii) A hierarchical campaign protocol: quality-diversity search, CMA-ES hill-climb on the top genome, then longer higher-resolution replay and a three-run Richardson resolution ladder (RC/RM/RF = coarse/medium/fine at fixed domain, varying grid spacing $h = L/N$) only on the refined champion (Sec. V).
- (iii) Explicit numerical limits of the surviving candidates: dispersing exotic matter and a transient ignition-phase controller—not passenger transport (Sec. VI).

II. MATTER-FIRST PIPELINE

A. Two-code toolchain

Two open numerical-relativity codes carry the loop. **GRTresna** [16] is a constraint solver: given a matter configuration it returns constraint-satisfying initial data ($\gamma_{ij}, K_{ij}, \phi, \Pi$) on the initial slice, so that the Hamiltonian and momentum constraints hold before any evolution. **GRTeclyn** [6, 7] is a GPU-accelerated, AMReX-based [17] code that evolves those data forward with the CCZ4 formulation [18] of the Einstein equations under

moving-puncture gauge [19, 20]. Neither code is told what metric to produce; the geometry is whatever the proposed matter sources.

B. Search Loop

One evaluation maps a parameter vector θ to a scalar fitness: $\theta \rightarrow \text{initial data} \rightarrow \text{GRTEclyn} \rightarrow \text{diagnostics} \rightarrow S(\theta)$ (Fig. 4). The optimizer proposes θ , **GRTresna** solves the constraints, **GRTEclyn** evolves the spacetime on GPU, and post-processing probes reduce the evolution to the fitness $S(\theta)$ that steers the next proposal.

The matter is built from complex scalar fields with canonical Klein–Gordon dynamics (two independent fields in the production model, Sec. II H); each field ϕ carries a conjugate momentum Π and potential $V(\phi)$. In the 3+1 split, a complex scalar field contributes an energy density ρ and a momentum density S_i that act as the sources in the Hamiltonian and momentum constraints,

$$\rho = \frac{1}{2}\psi^{-4}|\tilde{\nabla}\phi|^2 + \frac{1}{2}|\Pi|^2 + V(\phi), \quad S_i = -\text{Re}(\Pi^*\partial_i\phi), \quad (1)$$

where ψ is the conformal factor and $\tilde{\nabla}$ the conformal spatial gradient. The three terms in ρ are, in order, the gradient (spatial shape), kinetic (time variation), and potential (rest) energy of the field; S_i is nonzero only when the field oscillates in time ($\Pi \neq 0$) with a spatial gradient, which is what lets a moving scalar lump carry momentum and drive a shift. Equation (1) is therefore the precise statement of “the matter tells the geometry what to be”: θ sets ϕ and Π , and **GRTresna** returns the (γ_{ij}, K_{ij}) consistent with these sources.

The word “canonical” describes the Klein–Gordon equation of motion, not the sign with which every initial lump gravitates. For each lump the optimizer stores $x_{\text{ex},k} \in [0, 1]$, but configuration construction rounds it to $e_k = \text{round}(x_{\text{ex},k})$ and sets $s_k = (-1)^{e_k}$. Thus there is no physical fractional interpolation. In the **GRTresna** constraint solve the signed independent-lump approximation is

$$\rho_{\text{ID}} = \sum_k s_k \rho_k, \quad (S_i)_{\text{ID}} = - \sum_k s_k \sum_{A=1}^2 \Pi_{A,k} \partial_i \phi_{A,k}, \quad (2)$$

with inter-lump stress-tensor cross terms omitted. The profiles themselves are then superposed into two painted complex fields—one canonical, one phantom—by the bicomplex model. Candidate 146 rounds to $s_k = (+1, -1, -1, +1, -1)$: two canonical and three phantom sources (Sec. II H).

C. Q-ball building blocks

The matter is built from Q-balls [21, 22]—localized, non-topological solitons of a complex scalar field whose conserved Noether charge provides self-binding against

dispersal. Simpler ansatz families lacked this binding and disperse too rapidly (Sec. VI A). We choose a *complex* rather than real scalar because only a complex field has the phase symmetry that supports a stable charged lump, and its harmonic phase rotation $\phi \propto e^{-i\omega t}$ makes $\Pi = \dot{\phi} \neq 0$, so by Eq. (1) the lump carries a momentum density S_i and can source a nontrivial shift—the ingredient a transport channel needs. Each lump starts from a spherically symmetric equilibrium profile obtained by solving the radial ODE for a sextic potential,

$$V(|\Phi|) = \frac{1}{2}m^2|\Phi|^2 - \frac{1}{4}\lambda|\Phi|^4 + \frac{1}{6}\mu|\Phi|^6, \quad (3)$$

with fixed parameters $m = 1$, $\lambda = 640$, $\mu = 85333$, $\omega = 0.8$. The attractive quartic ($\lambda > 0$) provides self-binding; the repulsive sextic ($\mu > 0$) is *required* for a stable three-dimensional Q-ball—a pure quartic is the critical case and either collapses or disperses. The specific values are fixed by three requirements rather than tuned for performance: $m = 1$ sets the mass (hence length) unit of the problem; λ and μ are then chosen so that the stable thick-wall window $\omega_{\text{min}} < \omega < m$ comfortably brackets the working frequency, where λ^2/μ sets $\omega_{\text{min}} \approx 0.316$; and $\omega = 0.8$ places the soliton in the thick-wall regime, so that the equilibrium radial profile falls to zero within a radius of ~ 7 code units and several lumps fit inside the evolution box. These microphysical constants are held fixed and are not part of the search.

D. Shell Parameterization

The search vector has 39 dimensions because the arrangement of the lumps, not the lump microphysics, is what is optimized. The shell contains five lumps—the maximum supported by the solver’s state layout—giving enough independent tilted orbits to produce multi-stream frame dragging while keeping the search space tractable. Each lump carries seven searchable placement parameters—orbital radius R_0 , phase, two tilt angles, radial velocity, angular (rotation) velocity, and a source-sign selector—giving $5 \times 7 = 35$ per-lump dimensions. Four collective variables set the shell’s breathing amplitude and frequency and its vertical oscillation amplitude and frequency, for $35 + 4 = 39$ in total (Fig. 2). Although sampled continuously, the sign variable is applied as the binary selector above; it lets the optimizer choose which initial-data lumps source positive or negative stress-energy.

E. Lump orbit parameterization

Each lump moves in its own tilted orbital plane. The support center of lump k is prescribed by

$$\begin{aligned} r_k(t) &= \max[r_{\text{min}}, R_{0,k} + v_{r,k}t + A_b \sin(\omega_b t)], \\ \varphi_k(t) &= \varphi_{0,k} + \omega_{\text{rot},k}t, \quad z_b(t) = A_z \sin(\omega_z t), \end{aligned} \quad (4)$$




FIG. 2. Search parameterization of the trajectory Q-ball shell: $5 \times 7 = 35$ per-lump placement parameters plus 4 collective modes give the 39-D genome θ . MAP-Elites stores one elite per cell of an 8×8 archive indexed by peak shortcut and FTL lifetime. Q-ball microphysics is held fixed.

followed by rotation into the laboratory frame through the two tilt angles (θ_k, ϕ_k) . Thus each lump spirals radially (drift $v_{r,k}$), rotates in its own plane (angular velocity $\omega_{rot,k}$; negative values are retrograde), and shares a collective radial breathing (A_b, ω_b) and vertical oscillation (A_z, ω_z) . The range of possible orbits includes retrograde and prograde configurations, fast and slow rotation, radially drifting spirals, and nearly static scaffolds.

F. Trajectory controller (PD pump)

During evolution, a proportional-derivative (PD) controller drives the two real components of each complex field toward a moving target profile at the prescribed support center:

$$\begin{aligned} \partial_t \Pi_A|_{PD} = & -w k_p (\phi_A - \phi_A^*) \\ & -w k_d (\Pi_A - \Pi_A^*), \quad A = 1, 2, \end{aligned} \quad (5)$$

with gains $k_p = 12$, $k_d = 7$, and w the product of a localized envelope and a constraint-sensitive governor. This is an external, non-Lagrangian source in the scalar momentum equations; no corresponding pump stress-energy is added to $T_{\mu\nu}$, so $\nabla_\mu T^{\mu\nu} = 0$ is not exact while the pump is active. The volume-integrated pump power $P = \int \alpha (\Pi \cdot S_\Pi) \sqrt{\gamma} dV$ recorded in `collapse_diagnostics.dat` integrates to $E_{\text{pump}} \simeq 1.6 \times 10^{-17}$ over the ignition window of the pump-on HQ run—consistent with floating-point noise relative to the matter amplitude ($|\phi|_{\text{max}} \sim 7 \times 10^{-2}$). The pump therefore acts as a soft trajectory guide rather than a large energy injector. The campaign default switches the pump off at $t_{\text{pump}} = 4$; scoring credits 4D geodesic transport only for $t_{\text{emit}} \geq t_{\text{pump}}$. A dedicated pump-free HQ twin with $t_{\text{pump}} = 0$ (Sec. VI) tests whether the shortcut survives with the PD term identically off from $t = 0$.

The controller should be read as a *numerical guide*, not as a physical part of the reported spacetime: the headline geodesic launches at $t_{\text{emit}} = 4 = t_{\text{pump}}$, so the metric it samples is freely evolved ($\nabla_\mu T^{\mu\nu} = 0$), and the pump-free twin (Sec. VI) reproduces the advantage with the PD term identically off from $t = 0$.

G. Constraint gate and source-model consistency

A post-load constraint gate rejects inconsistent initial data before any GPU time is spent: the solved data are loaded into `GRTEclyn` for a short verification integration, and the volume-weighted L^2 norms of the Hamiltonian and momentum constraints must both remain below 10^{-2} , otherwise the candidate is discarded and logged. Candidate 146 passes this gate with $\|\mathcal{H}\|_2 = 5.9 \times 10^{-3}$ and $\|\mathcal{M}\|_2 = 1.1 \times 10^{-4}$. The campaign uses the `grtresna_bicomplex_scalar` matter model (Sec. IIH), which preserves per-lump gravitational signs through both the constraint solve and evolution. The post-load gate, together with bounded CCZ4 constraint growth, establishes numerical control and source-model consistency.

H. Phantom sector: action and evolution equations

Because the exotic sign is central to the claim, we state exactly what the `grtresna_bicomplex_scalar` model integrates. Each complex scalar is canonical, with action

$$S[\Phi] = \int d^4x \sqrt{-g} (-g^{\mu\nu} \partial_\mu \Phi^* \partial_\nu \Phi - V(|\Phi|)), \quad (6)$$

and V the sextic potential of Eq. (3). Writing $\Pi = n^\mu \partial_\mu \Phi$ for the conjugate momentum along the normal n^μ , both the canonical field Φ_+ and the exotic field Φ_- obey the same 3+1 Klein-Gordon system,

$$\begin{aligned} \partial_t \Phi_\pm &= \alpha \Pi_\pm + \beta^i \partial_i \Phi_\pm, \\ \partial_t \Pi_\pm &= \beta^i \partial_i \Pi_\pm + \alpha \left(K \Pi_\pm + \gamma^{ij} D_i D_j \Phi_\pm - \frac{\partial V}{\partial \Phi_\pm^*} \right) \\ &\quad + \gamma^{ij} \partial_i \alpha \partial_j \Phi_\pm, \end{aligned} \quad (7)$$

with D_i the spatial covariant derivative and $\alpha, \beta^i, \gamma_{ij}$ the usual lapse, shift, and spatial metric. The two sectors are therefore governed by identical, right-sign (non-ghost) kinetic operators; there is no wrong-sign kinetic term and hence none of the gradient or ghost instabilities that afflict a genuine phantom Lagrangian $\mathcal{L} = +\partial_\mu \phi \partial^\mu \phi^* - V$.

The exotic character enters at exactly one place: the *gravitational coupling*. The two sectors add to the Einstein source with opposite sign,

$$G_{\mu\nu} = 8\pi (T_{\mu\nu}[\Phi_+] - T_{\mu\nu}[\Phi_-]), \quad (8)$$

so Φ_- is a negative-energy (ghost-gravitational) source while retaining canonical, well-posed field dynamics.

This sign flip is applied consistently to ρ , S_i , and S_{ij} in both the **GRtresna** constraint solve [Eq. (2)] and the **GRteclyn** stress tensor, which is what removes the initial-data / evolution mismatch of earlier single-field campaigns. We are explicit that Eq. (8) is equivalent to reversing the effective gravitational constant for the exotic sector ($G \rightarrow -G$ for Φ_-); it is a controlled model of the negative-energy source required to violate the null energy condition, not a claim that a fundamental phantom field theory has been made well posed. The practical payoff of this choice is that the late-time dispersion reported below (Sec. VI A) is *structural*—the finite lifetime of thick-wall Q-balls together with an unbound orbital arrangement—rather than a manifestation of a phantom kinetic instability, since the fields themselves never leave the stable Klein–Gordon sector.

I. Scoring: how every evaluation is rewarded

Every evaluation in the campaign is reduced to one fitness value $S(\theta)$ by the same fixed rule, so candidates are ranked consistently. Each named quantity below is a normalized diagnostic in $[0, 1]$ (penalties are negative); the score is their weighted sum,

$$S = 1000 f_{\text{geo}}^{\text{evol}} + 1000 f_{\text{geo}}^{\text{op}} + 600 P_{\text{ftl}} + 200 F_{\text{op}} + 200 C_{\text{curv}} + 5 G_{\text{nt}} + g_{\text{h}} \Sigma_{\text{health}} + 40 w_x X_{\text{exotic}} + 200 H_{\text{hor}}, \quad (9)$$

with the health block

$$\Sigma_{\text{health}} = 150 \text{ surv} + 60 E + 40 \text{ stab} + 30 \text{ cmv} + 10 \mathcal{H}_{\text{ok}} + 15 I, \quad (10)$$

gated by the nontriviality factor $g_{\text{h}} \in [0, 1]$ so that a trivial (flat) configuration earns no health reward. Here $f_{\text{geo}}^{\text{evol}}$ is the timing advantage of Eq. (13) measured on the time-evolving metric stack, and $f_{\text{geo}}^{\text{op}}$ is the same geodesic quantity evaluated on frozen slices (a diagnostic, not the transport certificate; Sec. IV). Table I lists what each term measures. Two design choices matter. First, transport dominates: the two geodesic terms carry the largest weights, so the score chiefly tracks genuine null-transport gain. Second, the persistence and operational terms $P_{\text{ftl}}, F_{\text{op}}$ are *dispersion-gated*—multiplied by the fraction of matter still confined—so a candidate cannot bank a coordinate superluminal channel after its lumps have flown apart. The exotic penalty $X_{\text{exotic}} < 0$ charges for energy-condition violation [23] at campaign weight w_x . This penalty is deliberately sub-dominant to transport: at the campaign value $w_x = 0.2$ and the graded floor $X_{\text{exotic}} \geq -1.6$, the exotic term reaches at most $40 w_x |X_{\text{exotic}}| \approx 13$ points, under 7% of the ≈ 200 points a 20% evolving shortcut earns ($1000 f_{\text{geo}}^{\text{evol}}$). The optimizer is thus permitted to spend exotic matter freely in exchange for genuine null-transport gain, which is why the archive populates the phantom-supported sector rather than being pinned to $X_{\text{exotic}} = 0$. Search-tier

TABLE I. Reward terms in the campaign objective [Eq. (9)], applied identically to every MAP-Elites evaluation.

symbol	rewards / penalizes	weight
$f_{\text{geo}}^{\text{evol}}$	4D evolving null shortcut	1000
$f_{\text{geo}}^{\text{op}}$	geodesic (frozen) shortcut	1000
P_{ftl}	FTL persistence (gated)	600
F_{op}	operational coord. FTL (gated)	200
C_{curv}	curvature activity	200
G_{nt}	nontrivial geometry	5
surv	structural survival	150 g_{h}
E	energy-condition health	60 g_{h}
stab	metric stability	40 g_{h}
cmv	co-moving stability	30 g_{h}
I	instability penalty	15 g_{h}
$\mathcal{H}_{\text{ok}}$	constraint health	10 g_{h}
X_{exotic}	exotic-matter penalty	40 w_x
H_{hor}	horizon-formation penalty	200

scores are values of S on a $t = 16$ evolution (Sec. II J); the only production S values appear in the pump-free comparison (Table VII).

For the CMA-ES champion (candidate 146) the search-tier score is $S = 852.15$ (Sec. III).

J. Search and validation stages

MAP-Elites and CMA-ES evaluate $\mathcal{O}(10^2)$ genomes, so the discovery loop uses a reduced GRteclyn evolution configuration for wall-clock speed. Campaign env variables **GRtresna_evolution_L_full**, **GRtresna_evolution_N_full**, **GRtresna_evolution_max_level**, and **GRtresna_evolution_regrid_threshold** set the search-tier AMReX inputs **max_level** and **regrid_threshold** together with the evolution box (L, N) and t_{stop} . Archive / CMA scores are therefore *search-tier* diagnostics, not continuum claims. Only the frozen champion is re-run at production resolution for validation; Table II contrasts the two grid tiers and Fig. 3 summarizes the campaign ladder.

- **Stage 1 — MAP-Elites discovery (fast search grid):** quality-diversity search at $L = 64$, $N = 128$ (grid spacing $h \equiv L/N = 0.5$), **max_level**= 1, **regrid_threshold**= 0.1, and $t_{\text{stop}} = 16$ (bicomplex matter; per-lump gravitational signs retained in evolution). Coarse AMR is intentional: it buys enough evaluations to fill the behavior archive. Produces a diverse archive and a provisional score leader.
- **Stage 2 — CMA-ES refinement (same search grid):** warm-start from the MAP-Elites

TABLE II. GRTEclyn evolution settings: search tier (MAP-Elites / CMA-ES) versus production validation (HQ / Richardson medium). Search uses a reduced grid for throughput; validation upgrades L , N , max_level , and t_{stop} . Campaign env: `GRTRESNA_EVOLUTION_{L,N}_FULL`, `_MAX_LEVEL`, `_REGRID_THRESHOLD`.

parameter	search	HQ / RM
L	64	128
N	128	256
$h = L/N$	0.5	0.5
max_level	1	3
regrid_threshold	0.1	0.02
t_{stop}	16	30




FIG. 3. Outer campaign ladder. MAP-Elites finds the diverse archive and champion (87); CMA-ES refines it to candidate 146; only the refined champion enters high-resolution replay and the resolution matrix.

champion (and optionally its nearest elites); local covariance-matrix adaptation at the *same* ($L, N, \text{max_level}, t_{\text{stop}}$) and scoring rule. Only the refined top genome is promoted further.

- **Stage 3 — production replay (HQ):** the frozen CMA-ES champion is re-solved and evolved at higher fidelity / longer t_{stop} —typically $L = 128$, $N = 256$, $\text{max_level} = 3$, $\text{regrid_threshold} = 0.02$, $t_{\text{stop}} = 30$ —with multiple emission times and multi-ray geodesic probes. Paper headline numbers come from this stage, not from MAP-Elites.
- **Stage 4 — resolution and domain matrix:** matched coarse/medium/fine resolution runs (RC/RM/RF at fixed $L = 128$, $N = 192/256/384$) test continuum stability at the production max_level ; matched-spacing domain runs (DS/DL at $h = 0.500$, $L = 96/160$) test boundary sensitivity on the same t_{stop} window.

III. SEARCH METHODOLOGY

A. MAP-Elites quality-diversity

MAP-Elites [12] is a quality-diversity search: instead of returning one champion, it maintains an 8×8 archive with at most one elite per behavior cell. A candidate displaces the occupant of its cell only if its score $S(\theta)$ [Eq. (9)] is higher. The two archive axes are peak shortcut strength and FTL lifetime, so the campaign can keep short-lived strong channels and longer weaker ones side

TABLE III. Two-stage search hand-off: MAP-Elites champion (87) $\rightarrow$ CMA-ES refined champion (146) $\rightarrow$ HQ validation (RM). Search-tier grid: $L=64$, $N=128$, $\text{max_level}=1$, $t=16$. HQ grid: Table II.

stage	S	rays	peak f_{geo}	gated fitness
MAP-Elites (87)	825.19	3/3	25.35%	0.396
CMA-ES (146)	852.15	3/3	20.12%	0.417
HQ RM (146)	—	5/5	20.21%	—

by side. The bicomplex campaign scheduled 200 evaluations: 148 completed GPU evolution, 50 failed the post-load gate, one failed the GRTresna admission gate, and one failed the constraint solve. Nine cells were occupied (14.06% coverage). All nontrivial elites lie in the highest lifetime row, while peak strength spans the archive. The best score rose from 690.1 to 825.19 and last improved at batch 10, when candidate 87 entered cell (7, 7) (Fig. 5). The campaign runs to a fixed budget (200 evaluations here), which is set so that it comfortably outlasts saturation: it is terminated once the archive-best score has plateaued and cell coverage has stopped growing. In this run both stall well before the budget is exhausted (no score improvement after batch 10), so the fixed budget and the saturation criterion coincide.

B. Archive outcome and CMA-ES refinement

The bicomplex MAP-Elites winner (candidate 87, $S = 825.19$) combines a search-tier evolving shortcut of $f_{\text{geo}} = 25.3\%$ with dispersion-gated fitness 0.396, constraint-growth score 0.983, and confinement retention 0.396. CMA-ES warm-started from this genome and ran 150 evaluations (144 GPU-ok) at the same search grid, raising the score to 852.15 (candidate 146, $\Delta S \approx +27$; Fig. 6). The improvement is not a bigger raw shortcut—the search-tier f_{geo} actually *decreased* from 25.3% to 20.1%—but a better confinement-gated persistence (0.396 $\rightarrow$ 0.417). CMA-ES converged to a genome with nearly static lump placement: very slow angular velocities and small breathing, producing a quasi-static scaffold that keeps matter confined longer (Fig. 8). Table III summarizes the two-stage hand-off and its promotion to HQ validation.

Search-tier confinement already falls to $\sim 40\%$ by $t = 16$, so longer production runs are expected to stress structural persistence even if the evolving null shortcut remains positive.




FIG. 4. Inner evaluation loop shared by MAP-Elites and CMA-ES. Failed constraint solves and post-load checks never enter GPU evolution; successful evaluations compete for archive cells (MAP-Elites) or update the CMA distribution.




FIG. 5. MAP-Elites running best score over 200 evaluations. The champion (candidate 87) enters cell (7, 7) at batch 10 with $S = 825.19$.




FIG. 6. CMA-ES refinement progress over 150 evaluations, warm-started from MAP-Elites champion 87. Dashed line: MAP-Elites best (825.19). The running best exceeds the baseline early and reaches 852.15 at candidate 146.

IV. DIAGNOSTICS AND PROBE VALIDATION

A. The 4D evolving null-geodesic trace

The transport claim rests on tracing light rays through the simulated spacetime, not on reading coordinate speeds off a slice. A coordinate operational shortcut F_{op} [Eq. (12)] can flag tilted light cones, but it is partly gauge dependent and can be inflated by a coordinate choice. The authoritative probe therefore integrates the null-geodesic equations

$$\dot{x}^\mu = g^{\mu\nu} k_\nu, \quad \dot{k}_\mu = -\frac{1}{2}(\partial_\mu g^{\alpha\beta}) k_\alpha k_\beta, \quad (11)$$

for a photon with position x^μ and momentum k_μ . “4D” means the ray is not confined to one snapshot: it is pushed through a stack of stored time slices, so at each step the metric $g_{\mu\nu}(t, \mathbf{x})$ used is the one that actually exists when the photon is at that point. This distinguishes

it from a *frozen* trace, which wrongly assumes a single slice persists for the whole flight.

The construction is as follows. A ray is emitted at time t_{emit} from a fixed source point toward a fixed receiver. The quantities that enter the shortcut formulas are *transit durations*, not absolute grid arrival times: t_B^{curved} (respectively t_{min}) is the coordinate time elapsed along the curved path after emission, and t_B^{flat} (respectively t_{flat}) is the matched flat-space flight duration between the same endpoints. Absolute arrival on the grid is $t_{\text{arrival}} = t_{\text{emit}} + t_B^{\text{curved}}$. The fractional advantage is

$$F_{\text{op}} = \frac{t_B^{\text{flat}} - t_B^{\text{curved}}}{t_B^{\text{flat}}}, \quad (12)$$

$$f_{\text{geo}} = \max\left(0, \frac{t_{\text{flat}} - t_{\text{min}}}{t_{\text{flat}}}\right), \quad (13)$$

so $f_{\text{geo}} = 0$ means no shortcut and $f_{\text{geo}} = 0.20$ means the curved transit is 20% shorter than the flat comparison. A launch is accepted only if its rays stay physical: the relative drift in the null Hamiltonian $g^{\mu\nu} k_\mu k_\nu = 0$ must remain below 10^{-2} , which rejects rays that have accumulated interpolation error. The simulation itself runs to $t = 30$, but that is the evolution horizon of the metric stack, not the emission schedule. A ray launched at t_{emit} still needs a long remaining segment of evolved geometry to reach the far endpoint (flat transit is $t_{\text{flat}} = 14.4$), so late launches near $t = 30$ cannot finish a fair comparison. We therefore sweep seven early launch times, $t_{\text{emit}} = 0, 2, \dots, 12$ (five rays each), and report the largest accepted f_{geo} —the emission time at which the channel is most open. Across the pump-on validation matrix the peak sits at $t_{\text{emit}} = 4$ (just after the PD pump switches off) and has fallen to zero by $t_{\text{emit}} = 12$, so for those runs the window brackets the full life of the channel. (In the pump-free twin the optimum instead sits at the edge of the same window; Sec. VI.)

B. Null-constraint accuracy along the flight path

The 10^{-2} acceptance gate is a conservative safety bound, not the achieved accuracy. Two mechanisms keep the ray far tighter. First, whenever the absolute null Hamiltonian exceeds 10^{-6} the momentum is re-projected onto the local null cone before the next Runge-Kutta step, so interpolation error cannot accumulate secularly.




FIG. 7. Normalized null constraint $k_\mu k^\mu / (k^t)^2$ along the headline candidate-146 ray ($t_{\text{emit}} = 4$, RM) from emitter A ($x = 56.8$) to receiver B ($x = 71.2$). The ray stays below the 10^{-2} acceptance gate (dashed) throughout: the median is 3.0×10^{-4} , and the brief peak of 4.1×10^{-3} in the strong-field core remains a factor ~ 2.4 under the gate.

Second, the normalized null constraint stays small between projections. Figure 7 plots the gauge-independent ratio $k_\mu k^\mu / (k^t)^2$ along the headline candidate-146 ray ($t_{\text{emit}} = 4$, RM) as it crosses the matter region. The median value, 3.0×10^{-4} , sits a factor ~ 30 below the 10^{-2} gate; the peak, 4.1×10^{-3} , is reached briefly in the strong-field core and remains a factor ~ 2.4 below the gate. The ray therefore stays null to the stated tolerance throughout the flight and does not drift toward a timelike or spacelike trajectory. The complementary relative-drift diagnostic recorded by the probe, $h_{\text{rel}}^{\text{max}} = 2.19 \times 10^{-3}$ in RM, tells the same story.

C. Probe controls and calibration

Three controls bracket the probe's dynamic range before any candidate is evaluated.

- **Minkowski (flat):** $f_{\text{geo}} = 0$ by construction; verifies that the tracer introduces no spurious shortcut on trivial data.
- **Prescribed Alcubierre [3]:** a moving-bubble metric ($v_s = 0.5$, bubble radius 2, wall width 2) is built analytically and fed to the evolving tracer, yielding $f_{\text{geo}} \approx 32\%$ with 5/5 rays reaching. This confirms that the 4D probe recovers a known coordinate shortcut.
- **Negative control (eval086):** an earlier campaign candidate whose frozen-slice diagnostic shows $f_{\text{geo}} \approx 5.8\%$, but whose 4D evolving HQ trace collapses to 0% (0/5 rays). This false positive motivated the design rule that the evolving trace is authoritative: frozen peaks alone do not certify transport.

The positive campaign control is the CMA-ES champion (candidate 146): every accepted launch in all three

resolution-matrix runs has 5/5 rays reaching and passes the null-Hamiltonian gate.

V. RESULTS: CANDIDATE-146 WARP PRECURSOR

A. Provenance and genome

Candidate 146 is the CMA-ES-refined champion from the bicomplex campaign (Sec. III). Its provenance freezes the exact 39-dimensional genome, source checksums, solver parameters, and source commit, using the `gtrtesna_bicomplex_scalar` matter model. The finest run (RF) additionally stored field maps, from which the snapshots of Fig. 10 are drawn.

B. Matter dynamics

Candidate 146 uses five lumps with source signs $s_k = (+1, -1, -1, +1, -1)$: two canonical and three phantom (exotic). The collective breathing is small ($A_b = 0.083$, $\omega_b = 0.429$), the vertical oscillation is moderate ($A_z = 2.32$, $\omega_z = 0.016$), and the initial radii are $R_0 = (3.69, 5.50, 2.83, 5.59, 2.61)$. All angular velocities are negative, $\omega_{\text{rot}} = (-0.0017, -0.0007, -0.0031, -0.0009, -0.0049)$, so all five trajectories are retrograde but extremely slow. Radial drifts are small [$v_r = (-0.060, -0.010, 0.001, -0.005, -0.006)$]; the slow rotation and drift—roughly an order of magnitude quieter than MAP-Elites champion 87—is what CMA-ES converged to in order to maximize confinement-gated persistence.

Figure 8 shows these prescribed support-center trajectories [Eq. (4)] at four times. The Q-ball fields are initially boosted along the corresponding motion. The PD trajectory controller (Sec. II F) drives the lumps during the ignition phase ($t < 4$) and then switches off. The bounded constraints ($\|\mathcal{H}\|_2 \leq 4.73 \times 10^{-3}$, $\|\mathcal{M}\|_2 \leq 9.16 \times 10^{-4}$ in RM) show that the evolution remains numerically controlled. Figure 9 complements the magnified three-dimensional view with an unscaled top-down (x - y) projection at grid scale, showing that the shell spans only a few code units across the null-probe corridor. The support centers are the intended dynamical scaffold, not tracked density isosurfaces. The field follows them imperfectly: confinement falls from 72% to 11% and rms spread grows by a factor 3.27 by $t = 30$. Candidate 146 is therefore a transient, initially driven configuration, not a bound five-body soliton.

C. Transport mechanism

The shortcut arises from the interplay of matter momentum and geometry. Each Q-ball lump oscillates with




FIG. 8. Three-dimensional dynamics encoded by the frozen candidate-146 genome. Panels show the five prescribed PD-controller centers at $t = 0, 4, 16$, and 30 . Filled circles are the centers at the stated times. To make the small motion legible, each local trajectory glyph covers four code-time units before and after the snapshot (where available), with its displacement about the filled center magnified by $4\times$; the hollow circle and arrow mark the start and end of that window. Colors and labels identify the lumps (+ canonical, - phantom), and the dashed line marks the x -directed null-probe axis. The magnification is representational: the centers are plotted at their unscaled positions. CMA-ES converged to very slow angular velocities ($|\omega_{\text{rot}}| \lesssim 5 \times 10^{-3}$) and small breathing ($A_b = 0.083$), producing a quasi-static scaffold that maximizes confinement-gated persistence.




FIG. 9. Top-down (x - y) projection of the candidate-146 initial lump centers at true grid scale (box center $(x_0, y_0) = (64, 64)$; no magnification). Black circles are the two canonical (+ ρ) lumps, red squares the three phantom ($-\rho$) lumps [signs (+, -, -, +, -)]. The dashed line is the x -directed null-probe corridor from emitter A to receiver B ($|x - x_0| \leq 7.2$). Compared with Fig. 8, this panel shows the actual spatial extent of the shell: the five solitons occupy a region a few code units across, comparable to the corridor half-length.

$\Pi \neq 0$ [Eq. (1)], sourcing a momentum density S_i that enters the momentum constraint and generates a nontrivial shift β^i . The shift tilts the local light cones in the direction of the matter current; null rays propagating through this region accumulate a coordinate-time advance relative to the flat reference path (Fig. 14). The finest-run field maps in Fig. 10 make this concrete: the phantom and

TABLE IV. Evolving null-geodesic transport certificate for candidate 146 in the headline RM run ($L = 128$, $N = 256$, $t = 30$).

quantity	value
peak 4D f_{geo}	20.21%
frozen-peak f_{geo}	29.45%
rays reached	5/5
t_{emit}	4.00
t_{arrival}	15.49
$t_{\text{min}} = t_{\text{arrival}} - t_{\text{emit}}$	11.49
$t_{\text{transit}}^{\text{flat}}$	14.40
$h_{\text{quality_ok}}$	true ($h_{\text{rel}}^{\text{max}} = 2.19 \times 10^{-3}$)
search-tier peak (CMA-ES)	20.12% (3/3 rays)

canonical lumps (ρ_{req} , top row) source a bipolar shift β^x (bottom row) that is strongest at $t = 4$ and decays as the matter disperses. The nearly static, retrograde scaffold selected by CMA-ES (Sec. III) holds this shift structure in place while the matter remains confined.

The peak f_{geo} occurs at $t_{\text{emit}} = 4$ (Fig. 11); at later emission times the channel weakens as matter disperses. Lower-scoring archive configurations failed for the same reason: their supporting matter dispersed (Sec. VIA).

D. 4D falsification certificate

The headline claim is operational: in the RM run, null rays between the fixed endpoints arrive earlier than in flat space. Table IV lists the quantities reported by the evolving-geodesic probe.

Explicitly, $(t_{\text{flat}} - t_{\text{min}})/t_{\text{flat}} = (14.40 - 11.49)/14.40 = 20.21\%$, matching Eq. (13). The production HQ value




FIG. 10. Evolving dynamics of candidate 146 in the finest run (RF, $N = 384$), mid-plane ($z = z_0$) slices at $t = 0, 4, 16, 30$. *Top row*: required energy density ρ_{req} , with blue the canonical ($+\rho$) lumps and red the phantom ($-\rho$) lumps—a direct view of the negative-energy sector of Sec. II H. *Bottom row*: the shift component β^x along the null-probe axis, whose sign tilts the local light cones and opens the transport channel. The channel is strongest just after pump-off ($t = 4$, the peak- f_{geo} launch); by $t = 30$ both the matter and the shift structure have spread and weakened, consistent with the confinement drop from 72% to 11% and the transient nature of the shortcut.




FIG. 11. Evolving null-geodesic emission sweeps ($t_{\text{emit}} = 0, 2, \dots, 12$) for all three resolution-matrix runs; the metric evolves to $t = 30$, but late launches lack enough remaining geometry for a complete transit. All runs peak at $t_{\text{emit}} = 4$, just after PD pump-off, and agree within 0.26 percentage points. Every launch has 5/5 rays reaching and passes the null-Hamiltonian quality gate; at $t_{\text{emit}} = 12$ the rays arrive no earlier than the flat baseline, so the clamped advantage [Eq. (13)] is zero.

agrees with the search-tier value to within 0.5% relative. The endpoint timing corresponds to an effective $c/(1 - f_{\text{geo}}) \simeq 1.25c$ relative to the matched flat path; nothing physical moves outside its local light cone (Fig. 14).

E. Endpoint gauge and proper-time invariance

The timing advantage of Eq. (13) compares two *coordinate* transit durations between the same fixed endpoints. In moving-puncture gauge this is only a physical statement if the endpoints behave as nearly stationary, nearly flat observers; a deep lapse well ($\alpha \ll 1$) at the endpoints could in principle manufacture a coordinate advance with no proper-time content. We therefore report the metric at the two probe endpoints. The emitter and receiver sit on the propagation axis at $A = (x, y, z) = (56.8, 64, 64)$ and $B = (71.2, 64, 64)$ in the $L = 128$ box centered on $(64, 64, 64)$ (a corridor of coordinate length $t_{\text{flat}} = 14.4$). Table V lists the lapse α and shift magnitude $|\beta^i|$ there through the emission window. At the headline launch $t_{\text{emit}} = 4$ the endpoints are close to flat and stationary: $\alpha_A = 0.96$, $\alpha_B = 0.99$, and $|\beta^i| \lesssim 3.5 \times 10^{-2}$. The lapse never enters the dangerous $\alpha \ll 1$ regime; at late times it drifts modestly *above* unity ($\alpha \lesssim 1.2$), which if anything makes the coordinate advantage conservative, since a stationary clock with $\alpha > 1$ accumulates proper time faster than coordinate time.

To remove the residual gauge dependence entirely, we re-express the timing advantage in the proper time of a static receiver at B , whose clock ticks at $d\tau = \sqrt{-g_{00}(B)} dt$. Integrating the measured $g_{00}(B, t)$ from t_{emit} to each arrival, the curved signal reaches B at proper time $\tau_{\text{curved}} = 12.61$, while the matched flat-rate signal (coordinate arrival $t_{\text{emit}} + t_{\text{flat}}$) reaches it at $\tau_{\text{flat}} = 15.66$,

TABLE V. Lapse α and shift magnitude $|\beta^i|$ at the fixed probe endpoints A (emitter) and B (receiver) in RM, sampled through the emission window. The endpoints remain near-stationary ($|\beta^i| \ll 1$) and out of any deep-lapse well ($\alpha \geq 0.95$ at launch). Source: `endpoint_gauge_rm.txt`.

t	α_A	$ \beta^i _A$	α_B	$ \beta^i _B$
0	1.000	0	1.000	0
4	0.960	4.5×10^{-3}	0.993	3.5×10^{-2}
8	0.946	1.7×10^{-2}	1.048	9.3×10^{-2}
12	1.179	1.2×10^{-2}	1.209	1.3×10^{-2}
15	1.172	6.9×10^{-2}	1.118	7.1×10^{-2}

giving a proper-time advantage

$$f_{\text{geo}}^\tau = \frac{\tau_{\text{flat}} - \tau_{\text{curved}}}{\tau_{\text{flat}}} = 19.5\%, \quad (14)$$

within 0.7 percentage points of the coordinate value 20.2%. The advantage is therefore not a lapse/shift artifact: it survives conversion to the physical clock of a stationary receiver.

A complementary check concerns proper *distance*. The negative-energy sources contract the spatial metric along the corridor, so the coordinate separation 14.4 overstates the proper length of the A–B line. Integrating $\int \sqrt{\gamma_{xx}} dx$ along the corridor in the stored RM metric stack gives a proper length of 13.85 at $t = 0$ and 13.60 at the stored slice nearest the headline launch ($t = 4.32$), and the value remains in $[13.08, 13.67]$ throughout the flight window. Part of the coordinate-level advantage therefore reflects genuine spatial contraction sourced by the exotic matter—itself a physical effect, but one worth separating from light-cone tilting. Under the most conservative renormalization, comparing the curved transit $t_{\text{min}} = 11.49$ against a flat transit across the *smallest* instantaneous proper length, $(13.08 - 11.49)/13.08 = 12.2\%$, the advantage remains strictly positive. We quote the coordinate-referenced $f_{\text{geo}} = 20.2\%$ as the headline number because Eq. (13) defines the flat baseline by the fixed coordinate endpoints, but both normalizations certify an early arrival.

F. Resolution and domain matrix

Validation comprises two complementary tests of candidate 146 (Table VI). At fixed domain $L = 128$, a three-run resolution ladder—coarse RC ($N = 192$, $h = 0.667$), medium RM ($N = 256$, $h = 0.500$), and fine RF ($N = 384$, $h = 0.333$)—probes continuum stability. At fixed grid spacing $h = 0.500$, a domain ladder—DS ($L = 96$, $N = 192$) and DL ($L = 160$, $N = 320$)—probes sensitivity to outer-boundary placement. The medium-resolution, intermediate-domain run RM is the reference configuration quoted elsewhere in this paper. All evolutions impose non-periodic Sommerfeld (radiative) outer

TABLE VI. Resolution and domain matrix for candidate 146. Every run uses three refinement levels above the base grid (`max_level`= 3, `regrid_threshold`= 0.02) and terminates at $t = 30$. Peak f_{geo} is the maximum evolving 4D null-geodesic timing advantage; frozen peak is the maximum of the frozen-slice f_{geo} time series. RC/RM/RF change h at fixed L ; DS/DL change L at fixed h .

run	L	N	h	peak f_{geo}	frozen peak	t_{emit}
RC	128	192	0.667	20.44%	29.15%	4
RM	128	256	0.500	20.21%	29.45%	4
RF	128	384	0.333	20.19%	29.54%	4
DS	96	192	0.500	20.19%	29.21%	4
DL	160	320	0.500	20.22%	29.44%	4

boundary conditions on all six faces, so the domain ladder isolates the effect of moving those radiative boundaries in or out.

On the resolution ladder the evolving timing advantage spans 20.19–20.44%. The fine–medium relative difference is 0.13% (computed from unrounded values), well below the 10% acceptance tolerance fixed for the ladder before it was run. The frozen-peak f_{geo} remains near $\sim 29.5\%$ at all three resolutions. On the domain ladder, evolving $f_{\text{geo}} = 20.19\%$ (DS) and 20.22% (DL) differ from RM by at most 0.12% relative; each run returns five of five completed rays with the null-Hamiltonian quality gate satisfied. The agreement indicates that outer-boundary reflections do not set the reported timing advantage on the interval $t \leq 30$. That conclusion is consistent with a causal estimate for the $L = 128$ box: in $1 + \log$ slicing, gauge modes with characteristic speeds up to $\sim \sqrt{2}c$ require $t \gtrsim L/(2\sqrt{2}) \approx 45$ for a one-way trip from the outer boundary to the central region.

For unequal spacings, a Richardson estimate assumes

$$Q(h) = Q_0 + Ch^p, \quad \frac{Q_c - Q_m}{Q_m - Q_f} = \frac{h_c^p - h_m^p}{h_m^p - h_f^p}. \quad (15)$$

The f_{geo} sequence 20.44% $\rightarrow$ 20.21% $\rightarrow$ 20.19% is monotonically decreasing and continuum-stable (fine–medium relative difference 0.13%), but the absolute gap is too small for a clean order extraction from f_{geo} alone. Applying Eq. (15) directly to the triplet with the actual spacings $(h_c, h_m, h_f) = (0.667, 0.500, 0.333)$ returns an unstable exponent ($p \approx 8.7$) whose value is set by the fine–medium difference of a single significant figure; the sequence is thus consistent with rapid convergence but does not pin down p on its own. We therefore measure the observed order on the constraint norms themselves (Fig. 12). Over the ignition window $1.5 \leq t \leq 5$, the Hamiltonian residual converges at median observed order $p_{\mathcal{H}} \approx 3.3$ and the momentum residual at $p_{\mathcal{M}} \approx 2.8$, consistent with the high-order finite-difference scheme. At late times the three resolutions already sit on top of each other to within a few percent, so the order estimator becomes noise-dominated—as expected once the contin-




FIG. 12. Hamiltonian constraint norm across the Richardson ladder. The three resolutions overlie closely; the observed order extracted from Eq. (15) in the ignition window is $p_H \approx 3.3$ (median).




FIG. 13. Constraint norms in the headline RM run. The complete 3000-step diagnostic is plotted from a decimated table for typesetting only; maxima are evaluated from the full data.

uum residual is reached. We accept the 20.2% shortcut as continuum-stable without quoting a formal extrapolated value.

G. Evolution constraints

For RM, the volume-weighted norms remain below $\|\mathcal{H}\|_2 = 4.73 \times 10^{-3}$ and $\|\mathcal{M}\|_2 = 9.16 \times 10^{-4}$ through $t = 30$ (Fig. 13). These are comparable to the search-tier values and remain bounded throughout the evolution. The constraint norms peak during the early driven phase ($t \lesssim 4$) and then plateau or slowly grow during the free evolution, consistent with the dispersing matter producing less active refinement at late times.

TABLE VII. Matched HQ comparison of candidate 146 with the PD controller on through $t = 4$ (RM) and disabled from $t = 0$ (PF-RM). Both runs use $L = 128$, $N = 256$, and $t = 30$. Evolving f_{geo} is the 4D null-geodesic timing advantage; E_{pump} is the integrated pump work recorded in `collapse_diagnostics.dat`; S is the final `general_ftl` score.

metric	RM ($t_{\text{pump}} = 4$)	PF-RM ($t_{\text{pump}} = 0$)
evolving f_{geo}	20.21%	20.46%
rays / h quality	5/5, pass	5/5, pass
best t_{emit}	4	12
E_{pump}	1.6×10^{-17}	0
confined ($t=0 \rightarrow 30$)	72% $\rightarrow$ 11%	72% $\rightarrow$ 2%
rms spread ratio	3.27	3.07
final score S	126.3	47.0

H. Frozen versus evolving shortcut, and confinement

RM’s frozen-peak f_{geo} reaches 29.45%, compared with the evolving 4D sweep peak of 20.21%. The ~ 9 -percentage-point gap is physical: the frozen trace assumes that one snapshot persists for the entire ray flight, whereas the 4D trace samples the actual changing metric and therefore “sees” the channel open and then weaken during propagation. Only the latter is used for the transport claim. The frozen peak is stable across all three resolutions (29.1–29.5%), confirming that the metric snapshot quality is not resolution-sensitive.

The evolving f_{geo} curve declines steadily at later emission times; successful ray transport coexists with the substantial dispersal described in Sec. V.

I. Evolution with the PD controller disabled

Because the PD trajectory controller (Sec. IIF) is an external, non-Lagrangian source, we repeat the HQ evolution of candidate 146 with the controller disabled from $t = 0$ ($t_{\text{pump}} = 0$), holding all other settings fixed ($L = 128$, $N = 256$, three refinement levels, $t = 30$). Table VII compares this pump-free integration (PF-RM) with the default pump-on twin (RM; $t_{\text{pump}} = 4$).

With the PD term off, the evolving timing advantage remains $f_{\text{geo}} = 20.46\%$ —within 1.3% relative of the pump-on value—with all five rays reaching and the null-Hamiltonian quality gate satisfied. The integrated pump work vanishes identically, so the advantage cannot be ascribed to external PD energy injection. The emission optimum shifts from $t_{\text{emit}} = 4$ to $t_{\text{emit}} = 12$, indicating a slower opening of the null channel in the absence of the controller, while the peak amplitude is essentially unchanged. Because $t_{\text{emit}} = 12$ is the last point of the emission sweep, the pump-free value is strictly a lower bound on that run’s peak advantage; the channel may




FIG. 14. Spacetime diagram of the null-geodesic shortcut. A photon emitted at A propagates toward receiver B. In flat space (dashed) the transit takes t_{flat} . Inside the Q-ball matter region (shaded band) the geometry bends null paths so that the ray through the curved spacetime (solid) arrives at B in $t_{\text{min}} < t_{\text{flat}}$. The bracket on the right marks the $\sim 20\%$ coordinate-time saving. The ray never leaves its local light cone.

still be opening when the sweep window closes.

The principal effect of the brief ignition-phase controller is therefore on matter retention rather than on f_{geo} . At $t = 30$ the pump-on run still confines $\sim 11\%$ of the matter, compared with $\sim 2\%$ when the controller is off; the rms spread ratios remain comparable (~ 3). The composite score decreases from $S = 126.3$ (RM) to $S = 47.0$ (PF-RM) for the same reason: confinement- and horizon-gated terms penalize the more rapidly dispersing pump-free evolution, while the geodesic timing advantage itself is essentially unchanged. A short PD drive in the default campaign thus delays dispersal of the supporting lumps. In both integrations the observed shortcut remains a transient precursor—a null-timing advantage coexisting with dispersing exotic matter.

VI. DISCUSSION

A. The dispersion problem

Candidate 146 should be read against the many configurations that did not work. Across roughly a dozen ansatz families and several full MAP-Elites campaigns we repeatedly found the same failure mode: the supporting matter disperses and the transport channel closes with it. Spherical-harmonic shells yielded a single weak hit in ~ 200 evaluations. Free massive scalar packets have no bound state and blow apart almost immediately. Bicomplex phantom lumps *without* Q-ball self-binding (an earlier, pre-soliton family) opened a gauge-invariant shortcut but lost $\sim 98\%$ of their structure by $t = 16$ (structural persistence ≈ 0.02). Boosted and stiffened Q-ball variants improved early confinement but still sheared apart

once placed on an orbit. The dispersion gate in the score (Sec. III) prevents the optimizer from exploiting channels whose matter has already dispersed. CMA-ES improved confinement-gated persistence (Sec. III), but candidate 146 retains only $\sim 11\%$ by $t = 30$. In summary, a *transient* shortcut is reproducible and robust, while a *long-lived*, well-confined channel remains an open problem.

B. Canonical-only (positive-energy) search bound

Reported positive-energy warp solitons [14, 15] and their critique [11] make the positive-energy sector the natural next target for an automated search. To test whether the same five-lump bicomplex Q-ball genome can produce a gauge-invariant timing advantage without phantom sources, we reran the MAP-Elites campaign of Sec. III with every per-lump exotic amplitude pinned to zero, retaining the canonical Klein–Gordon sector, retrograde-orbit bias, and post-load constraint gates. Over 200 evaluations (105 GPU-accepted; 95 rejected at the post-load Hamiltonian L^2 gate), the archive champion reached only $S \approx 47.4$ (eval. 130), and the evolving geodesic fraction satisfied

$$\max_{\theta \in \mathcal{A}_{\text{can}}} f_{\text{geo}}^{\text{evol}}(\theta) = 0 \quad (N_{\text{ok}} = 105). \quad (16)$$

Several accepted genomes developed superluminal local coordinate signal speeds ($\max v_{\text{loc}} > 1$, a gauge-dependent diagnostic) or early apparent horizons, but none completed a null-timing advantage under the evolving-geodesic certificate. Within this ansatz and search budget, a completed $f_{\text{geo}} > 0$ therefore appears to require the phantom sector used by candidate 146; the positive-energy search remains open for broader matter families and longer t_{stop} .

C. Limitations of the matter-first approach

The matter-first approach (Sec. I) is complementary to geometry-first [5]: every candidate is constraint-satisfying by construction, but the search is confined to the chosen matter model. Three specific limitations follow.

1. **Ansatz dependence.** The pipeline discovers only what the parameterized matter family can express. The 39-dimensional bicomplex Q-ball genome covers orbital arrangements of five solitons but cannot represent, e.g., anisotropic fluids, higher-spin fields, or engineered equation-of-state shells. Broadening the matter sector is the most direct route to new mechanism families.
2. **$\mathbb{R}^3$ spatial topology.** The GRTresna constraint solve and the GRTeclyn Cauchy evolution both operate on a topologically trivial spatial slice ($\mathbb{R}^3$

with the non-periodic Sommerfeld radiative boundaries described in Sec. V). Topological features that exist only in the full four-dimensional manifold—wormhole throats, handles, or other nontrivial $\mathbb{R}^4$ structures—cannot be represented on the initial slice and are therefore invisible to the search. Discovering such configurations would require either an extended initial-data formalism (e.g., puncture or excision methods for nontrivial topology) or a geometry-first seed that is subsequently evolved.

3. Ignition-phase controller. The PD pump (Sec. IIF) is a non-Lagrangian external source; results without it (Sec. VI) are comparable. A fully self-consistent confinement mechanism would remove the need for any external drive.

The two approaches are therefore complementary: geometry-first is better suited to exploring exotic topologies and unconstrained metric deformations, while matter-first provides constraint-satisfying, dynamically evolvable candidates that can be scored by a 4D transport test. A hybrid strategy—using geometry-first motifs as seeds for a matter-first refinement loop—is a natural next step.

D. Applications and next steps

The pipeline supports (i) inverse design of null transport under energy-condition penalties, (ii) QD atlases of mechanism families, and (iii) Pareto maps of shortcut versus exotic cost and persistence. The canonical-only MAP-Elites bound of Sec. VIB already constrains the all- $+\rho$ slice of the present ansatz. Immediate priorities include longer t_{stop} integrations to track late-time dispersal and broader positive-energy matter families. Longer-term goals include non-dispersing matter models and confinement world-tube objectives.

E. Laboratory Casimir bounds on null-timing shortcuts

The exotic negative energy supporting candidate 146 is a continuum phantom sector, not a laboratory vacuum stress. For ideal parallel plates of separation a , the Casimir energy density [24] is

$$\rho_{\text{cas}} = -\frac{\pi^2 \hbar c}{720 a^4}. \quad (17)$$

Under linearized semiclassical gravity, even an optimistic $a = 1 \text{ nm}$ gap in a 1 km waveguide yields $f_{\text{geo}}^{\text{bound}} \sim 10^{-50}$ —nearly 50 orders of magnitude below candidate 146’s 20.2%, and ~ 37 orders below the best atomic-clock resolution, consistent with quantum-inequality bounds [25]. Proposed amplification routes (dynamical Casimir effect [26], plasmonic metamaterials,

slow-light media) remain open modelling targets. The automated search discovers large f_{geo} only because it explores continuum exotic matter far outside laboratory stress-energy budgets.

F. Computational cost

All runs were performed on a single 8-GPU node (NVIDIA H100 80 GB HBM3). MAP-Elites and CMA-ES each dispatch one GRTeclyn evolution per GPU in parallel, with concurrent GRTresna constraint solves running on CPU (MPI, 8 ranks per solve). A 200-evaluation MAP-Elites campaign at search resolution ($N = 128$, $t_{\text{stop}} = 16$) completes in approximately two wall-clock hours; the 150-evaluation CMA-ES refinement takes a similar time. The production validation stage is sequential: the headline RM run ($N = 256$, $t_{\text{stop}} = 30$, three refinement levels) fits on a single GPU; the fine Richardson run RF ($N = 384$) required three GPUs due to memory constraints. Total wall-clock time from the start of MAP-Elites to the completion of the three-run resolution matrix was approximately ten hours.

G. Future applications of the pipeline

The matter-first search architecture—MAP-Elites exploration, CMA-ES refinement, and automated scoring on the evolved spacetime—is not specific to null-timing shortcuts. Any question of the form “which initial matter configuration optimizes geometric observable $\mathcal{O}$?” can be attacked by replacing the scoring function while reusing the genome representation, constraint solver, and evolution infrastructure.

Directional gravitational-wave generation. A parallel campaign (in preparation) replaces f_{geo} with a beaming-weighted Weyl-scalar power score, searching for Q-ball configurations that emit preferentially along one axis. Early runs produce $\sim 30\%$ directional beaming gain with mean Weyl power $P \sim 10^{-6}$ – 10^{-4} , demonstrating that the optimizer can steer toward anisotropic radiation from matter that is individually spherically symmetric.

Other candidate targets. (i) Gravitational-lensing amplification—optimizing a matter shell to maximize the magnification of a background source at a specified image plane; (ii) critical-collapse threshold mapping—locating the boundary between dispersal and black-hole formation in multi-lump parameter space; (iii) boson-star merger waveform exploration—seeding binary or multi-body initial data and scoring the resulting gravitational-wave energy or kick velocity. In each case the pipeline inherits the hard gates (Hamiltonian constraint, collapse diagnostics) that reject numerical artefacts, addressing the reward-hacking failure modes common to optimization in nonlinear PDE solvers.

VII. CONCLUSION

We presented an automated matter-first discovery pipeline in which MAP-Elites proposes scalar configurations, CMA-ES refines the archive champion, and 4D null geodesics through the evolved spacetime test whether light arrives earlier than in flat space. The refined champion (candidate 146, a bicomplex five-lump Q-ball shell) yields $f_{\text{geo}} = 20.2\%$ in the headline validation run, stable across a three-run resolution ladder ($\sim 0.1\%$ fine-medium difference) and a matched domain ladder ($\leq 0.12\%$; Table VI). The pump-free twin reproduces the advantage at 20.5% (Sec. VI). The shortcut is transient; a matched canonical-only search found no $f_{\text{geo}} > 0$ (Sec. VIB). Candidate 146 is therefore a reproducible, continuum-stable numerical warp precursor—not a functional warp drive—and semiclassical Casimir bounds (Sec. VIE) confirm that laboratory vacuum stresses fall ~ 50 orders of magnitude short of the required exotic budget.

CODE AND DATA AVAILABILITY

The complete campaign archive—run configurations, frozen genomes, raw diagnostics, and the analysis scripts that generate every figure and table in this paper—is available from the author upon reasonable request, and a

public release of the repository is in preparation. Within the archive, the project repository contains the campaign configuration and archive, candidate-146 frozen provenance (under `runs/grtresna_promote/sources/`), and the completed validation summaries (RC/RM/RF resolution ladder, DS/DL domain ladder, and the pump-free PF-RM twin: evolving-geodesic reports, constraint histories, and confinement timeseries). The canonical-only MAP-Elites archive `qball_traj_canonical_v1` (Sec. VIB) lives under `runs/grtresna_qd/`. Figure files live in `research/neuralspacetime/figures/`; the article directory holds the reduced numerical tables used by TikZ/PGFPlots. The Casimir bound script (Sec. VIE) lives under `research/neuralspacetime/casimir_null_timing_bounds.py`. MAP-Elites and CMA-ES trajectory files are under `runs/grtresna_qd/` and `runs/grtresna_cmaes/`.

ACKNOWLEDGMENTS

This work was supported by Gravity Frontiers (<https://www.gravityfrontiers.org/en>), which funded the underlying research and the development of the numerical-simulation software. The author extends sincere gratitude to Ilya Nachevsky for generously providing the high-performance computing resources essential to this work, and acknowledges the GRTL Collaboration for GRChombo, GRtresna, and GRTeclyn.

-
- [1] S. Alexander, B. Bradley, L. Gouskos, and C. Niu, “Autonomous Discovery of Particle Physics Theories from Experimental Data,” arXiv:2603.28935 [hep-ph] (2026).
 - [2] S. Qiu, Z. Cai, J. Wei, Z. Li, Y. Yin, Q.-H. Cao, C. Liu, M.-x. Luo, X.-B. Yuan, and H. X. Zhu, “An End-to-end Architecture for Collider Physics and Beyond,” arXiv:2603.14553 [hep-ph] (2026).
 - [3] M. Alcubierre, “The warp drive: hyper-fast travel within general relativity,” *Class. Quantum Grav.* **11**, L73 (1994).
 - [4] J. Natário, “Warp drive with zero expansion,” *Class. Quantum Grav.* **19**, 1157 (2002).
 - [5] C. Helmerich, J. Fuchs, A. Bobrick, L. Sellers, B. Melcher, and G. Martire, “Analyzing warp drive spacetimes with Warp Factory,” *Class. Quantum Grav.* **41**, 095009 (2024).
 - [6] GRTL Collaboration, **GRTeclyn**: a GPU-accelerated, AMReX-based numerical-relativity code, <https://github.com/GRTLCollaboration/GRTeclyn>.
 - [7] T. Andrade *et al.*, “GRChombo: An adaptable numerical relativity code for fundamental physics,” *J. Open Source Softw.* **6**, 3703 (2021).
 - [8] K. D. Olum, “Superluminal travel requires negative energies,” *Phys. Rev. Lett.* **81**, 3567 (1998).
 - [9] S. Gao and R. M. Wald, “Theorems on gravitational time delay and related issues,” *Class. Quantum Grav.* **17**, 4999 (2000).
 - [10] M. Visser, B. Bassett, and S. Liberati, “Superluminal censorship,” *Nucl. Phys. B Proc. Suppl.* **88**, 267 (2000).
 - [11] J. Santiago, S. Schuster, and M. Visser, “Generic warp drives violate the null energy condition,” *Phys. Rev. D* **105**, 064038 (2022).
 - [12] J.-B. Mouret and J. Clune, “Illuminating search spaces by mapping elites,” arXiv:1504.04909 (2015).
 - [13] N. Hansen and A. Ostermeier, “Completely derandomized self-adaptation in evolution strategies,” *Evol. Comput.* **9**, 159 (2001).
 - [14] E. W. Lentz, “Breaking the warp barrier: hyper-fast solitons in Einstein–Maxwell-plasma theory,” *Class. Quantum Grav.* **38**, 075015 (2021).
 - [15] S. D. B. Fell and L. Heisenberg, “Positive energy warp drive from hidden geometric structures,” *Class. Quantum Grav.* **38**, 155020 (2021).
 - [16] J. C. Aurrekoetxea, S. E. Brady, L. Aresté-Saló, J. Bamber, L. Chung-Jukko, K. Clough, E. de Jong, M. Elley, P. Figueras, T. Helfer, E. A. Lim, M. Radia, and A. Waeming, “GRtresna: An open-source code to solve the initial data constraints in numerical relativity,” *J. Open Source Softw.* (2026), arXiv:2501.13046 [gr-qc].
 - [17] W. Zhang *et al.*, “AMReX: a framework for block-structured adaptive mesh refinement,” *J. Open Source Softw.* **4**, 1370 (2019).
 - [18] D. Alic, C. Bona-Casas, C. Bona, L. Rezzolla, and C. Palenzuela, “Conformal and covariant formulation of the Z4 system with constraint-violation damping,” *Phys.*

- Rev. D **85**, 064040 (2012).
- [19] M. Campanelli, C. O. Lousto, P. Marronetti, and Y. Zlochower, “Accurate evolutions of orbiting black-hole binaries without excision,” *Phys. Rev. Lett.* **96**, 111101 (2006).
 - [20] J. G. Baker, J. Centrella, D.-I. Choi, M. Koppitz, and J. van Meter, “Gravitational-wave extraction from an inspiraling configuration of merging black holes,” *Phys. Rev. Lett.* **96**, 111102 (2006).
 - [21] S. Coleman, “Q-balls,” *Nucl. Phys. B* **262**, 263 (1985); erratum **269**, 744 (1986).
 - [22] T. D. Lee and Y. Pang, “Nontopological solitons,” *Phys. Rep.* **221**, 251 (1992).
 - [23] A. T. Le, “Observer-robust energy condition verification for warp drive spacetimes,” *arXiv:2602.18023 [gr-qc]* (2026).
 - [24] H. B. G. Casimir, “On the attraction between two perfectly conducting plates,” *Proc. Kon. Ned. Akad. Wetensch.* **51**, 793 (1948).
 - [25] L. H. Ford and T. A. Roman, “Averaged energy conditions and quantum inequalities,” *Phys. Rev. D* **51**, 4277 (1995).
 - [26] C. M. Wilson, G. Johansson, A. Pourkabirian, M. Simoen, J. R. Johansson, T. Duty, F. Nori, and P. Delsing, “Observation of the dynamical Casimir effect in a superconducting circuit,” *Nature* **479**, 376 (2011).